

UC Berkeley
Teaching Professor
Dan Garcia

CS61C

Great Ideas in Computer Architecture (a.k.a. Machine Structures)



UC Berkeley
Professor
Bora Nikolić

Introduction to the C Programming Language

Memory

Dynamic Memory Allocation (1/4)

- C has operator **sizeof()** which gives size in bytes (of type or variable)
- Assume size of objects can be misleading and is bad style, so use **sizeof(type)**
 - Many years ago an **int** was 16 bits, and programs were written with this assumption.
 - What is the size of integers now?
- “**sizeof**” knows the size of arrays:

```
int ar[3]; // Or:    int ar[] = {54, 47, 99}  
sizeof(ar) → 12
```

- ...as well for arrays whose size is determined at run-time:

```
int n = 3;  
int ar[n]; // Or: int ar[fun_that_returns_3()];  
sizeof(ar) → 12
```

Dynamic Memory Allocation (2/4)

- To allocate room for something new to point to, use `malloc()` (with the help of a typecast and `sizeof`):

```
ptr = (int *) malloc (sizeof(int));
```

- Now, `ptr` points to a space somewhere in memory of size `(sizeof(int))` in bytes.
- `(int *)` simply tells the compiler what will go into that space (called a typecast).
- `malloc` is almost never used for 1 var
- `ptr = (int *) malloc (n*sizeof(int));`
 - This allocates an array of `n` integers.

Dynamic Memory Allocation (3/4)

- Once `malloc()` is called, the memory location **contains garbage**, so don't use it until you've set its value.
- After dynamically allocating space, we must dynamically free it:
 - `free(ptr) ;`
- Use this command to clean up.
 - Even though the program **frees** all memory on **exit** (or when **main** returns), don't be lazy!
 - You never know when your **main** will get transformed into a subroutine!

Dynamic Memory Allocation (4/4)

- The following two things will cause your program to crash or behave strangely later on, and cause VERY VERY hard to figure out bugs:
 - `free()` ing the same piece of memory twice
 - calling `free()` on something you didn't get back from `malloc()`
- The runtime does not check for these mistakes
 - Memory allocation is so performance-critical that there just isn't time to do this
 - The usual result is that you corrupt the memory allocator's internal structure
 - You won't find out until much later on, in a totally unrelated part of your code!

Managing the Heap: `realloc(p, size)`

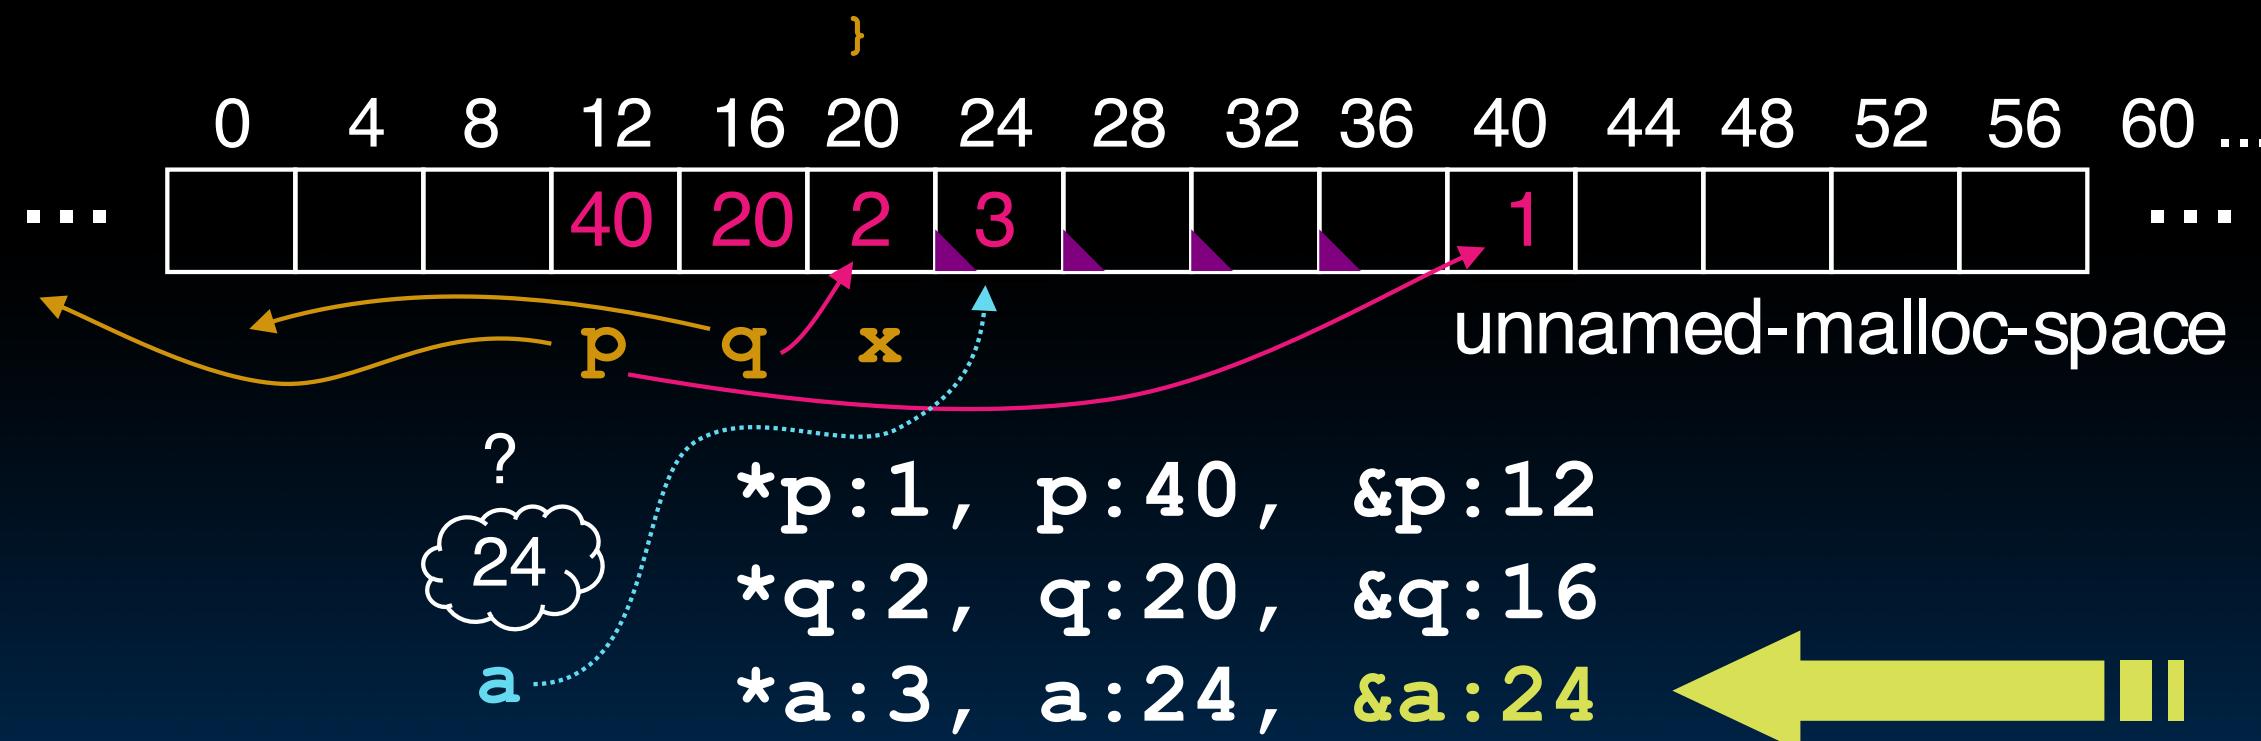
- Resize a previously allocated block at `p` to a new `size`
- If `p` is `NULL`, then `realloc` behaves like `malloc`
- If `size` is 0, then `realloc` behaves like `free`, deallocated the block from the heap
- Returns new address of the memory block; NOTE it is likely to have moved!

```
int *ip;  
ip = (int *) malloc(10*sizeof(int));  
/* always check for ip == NULL */  
... ... ...  
ip = (int *) realloc(ip,20*sizeof(int));  
/* always check NULL, contents of first 10  
elements retained */  
... ... ...  
realloc(ip,0); /* identical to free(ip) */
```

Arrays not implemented as you'd think

```
void foo() {  
    int *p, *q, x;  
    int a[4];  
    p = (int *)  
        malloc (sizeof(int));  
    q = &x;
```

```
*p = 1; // p[0] would also work here  
printf("*p:%u, p:%u, &p:%u\n", *p, p, &p);  
  
*q = 2; // q[0] would also work here  
printf("*q:%u, q:%u, &q:%u\n", *q, q, &q);  
  
*a = 3; // a[0] would also work here  
printf("*a:%u, a:%u, &a:%u\n", *a, a, &a);
```



K&R: “An array name is not a variable”

Mini-summary

- Pointers and arrays are **virtually same**
- C knows how to **increment pointers**
- C is an efficient language, with little protection
 - Array bounds not checked
 - Variables not automatically initialized
- Use handles to change pointers
- Dynamically allocated heap memory must be manually deallocated in C.
 - Use **malloc()** and **free()** to allocate and deallocate memory from heap.
- (Beware) The cost of efficiency is more overhead for the programmer.
 - “C gives you a lot of extra rope, don’t hang yourself with it!”



Linked List Example

Linked List Example

- Let's look at an example of using structures, pointers, `malloc()`, and `free()` to implement a linked list of strings.

```
struct Node {  
    char *value;  
    struct Node *next;  
};  
typedef struct Node *List;  
  
/* Create a new (empty) list */  
List ListNew(void)  
{ return NULL; }
```

Linked List Example

```
/* add a string to an existing list */
List list_add(List list, char *string)
{
    struct Node *node =
        (struct Node*) malloc(sizeof(struct Node));
    node->value =
        (char*) malloc(strlen(string) + 1);
    strcpy(node->value, string);
    node->next = list;
    return node;
}
```

node:



list



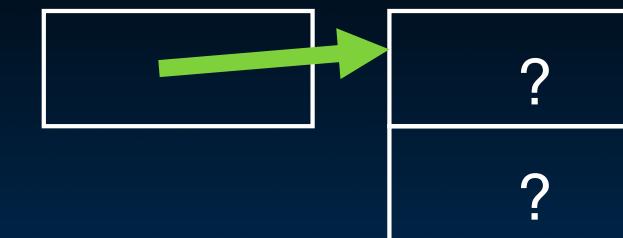
string:



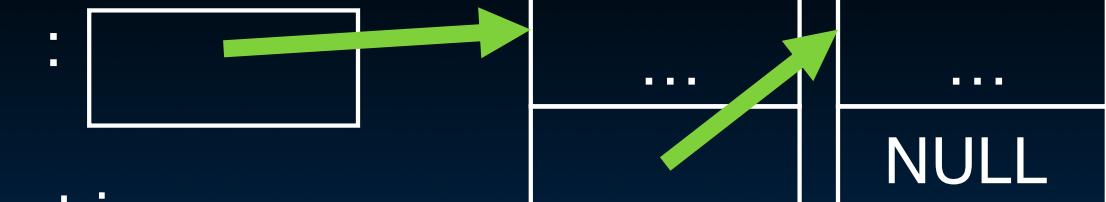
Linked List Example

```
/* add a string to an existing list */
List list_add(List list, char *string)
{
    struct Node *node =
        (struct Node*) malloc(sizeof(struct Node));
    node->value =
        (char*) malloc(strlen(string) + 1);
    strcpy(node->value, string);
    node->next = list;
    return node;
}
```

node:



list

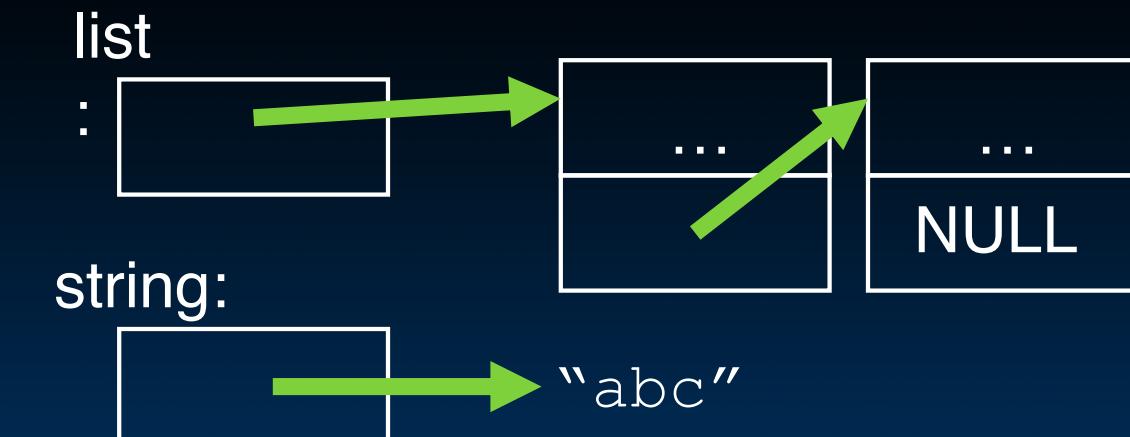
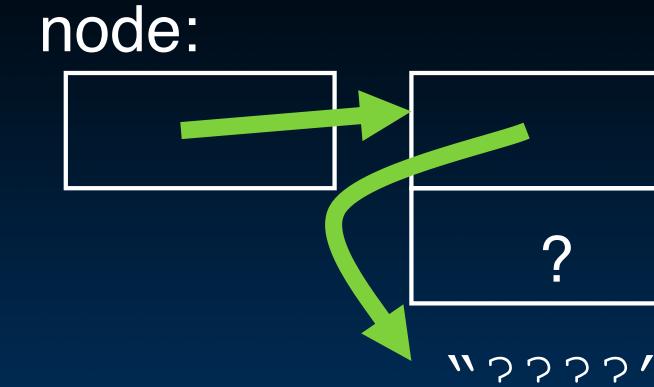


string:



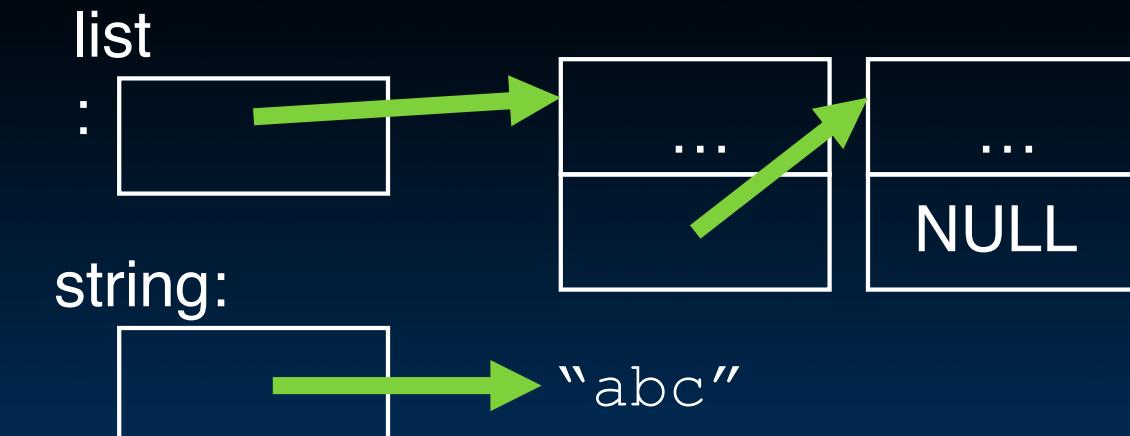
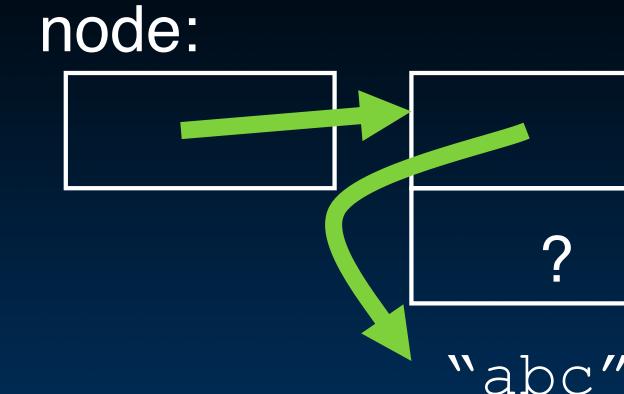
Linked List Example

```
/* add a string to an existing list */
List list_add(List list, char *string)
{
    struct Node *node =
        (struct Node*) malloc(sizeof(struct Node));
    node->value =
        (char*) malloc(strlen(string) + 1);
    strcpy(node->value, string);
    node->next = list;
    return node;
}
```



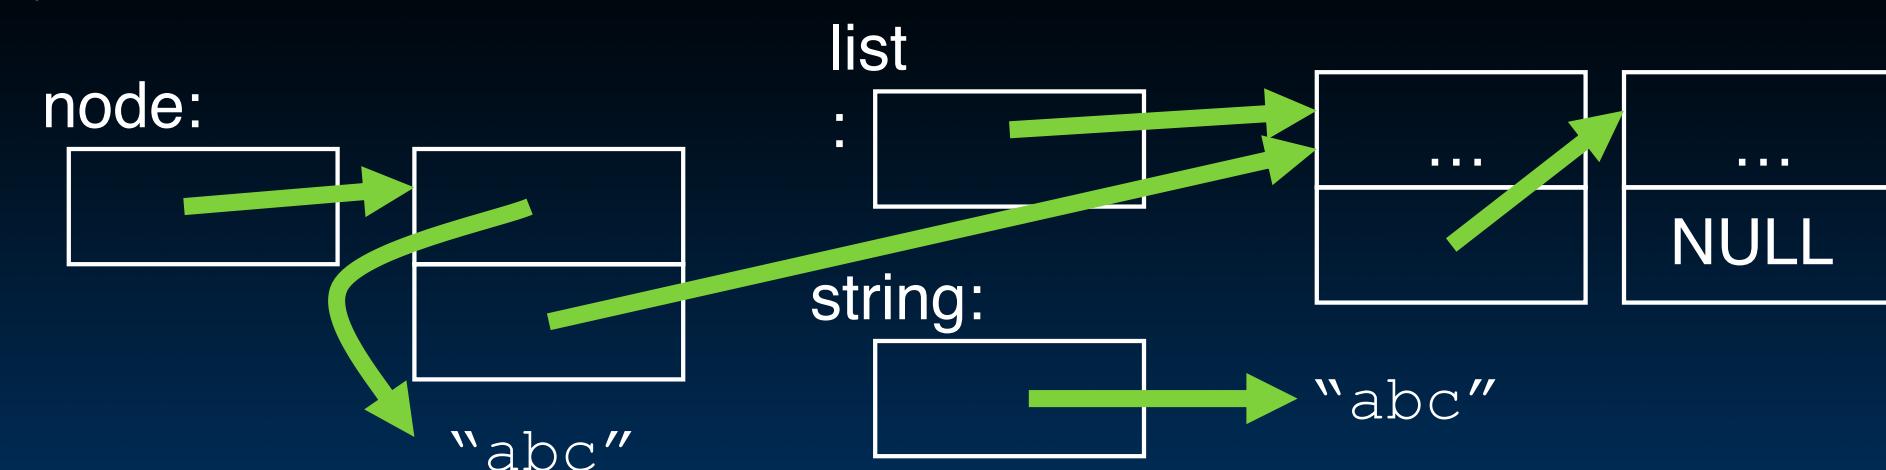
Linked List Example

```
/* add a string to an existing list */
List list_add(List list, char *string)
{
    struct Node *node =
        (struct Node*) malloc(sizeof(struct Node));
    node->value =
        (char*) malloc(strlen(string) + 1);
    strcpy(node->value, string);
    node->next = list;
    return node;
}
```



Linked List Example

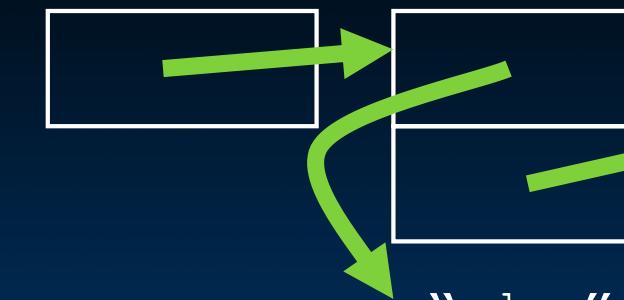
```
/* add a string to an existing list */
List list_add(List list, char *string)
{
    struct Node *node =
        (struct Node*) malloc(sizeof(struct Node));
    node->value =
        (char*) malloc(strlen(string) + 1);
    strcpy(node->value, string);
    node->next = list;
    return node;
}
```



Linked List Example

```
/* add a string to an existing list */
List list_add(List list, char *string)
{
    struct Node *node =
        (struct Node*) malloc(sizeof(struct Node));
    node->value =
        (char*) malloc(strlen(string) + 1);
    strcpy(node->value, string);
    node->next = list;
    return node;
}
```

node:



Memory Locations

Don't forget the globals!

- **What is stored?**
 - Structure declaration does not allocate memory
 - Variable declaration **does** allocate memory
- **So far we have talked about several different ways to allocate memory for data:**
 - Declaration of a local variable

```
int i; struct Node list; char *string; int ar[n];
```
 - “Dynamic” allocation at runtime by calling allocation function (alloc).

```
ptr = (struct Node *) malloc (sizeof(struct Node)*n);
```
- **One more possibility exists...**
 - Data declared outside of any procedure (i.e., before `main`).
 - Similar to #1 above, but has “global” scope.

```
int myGlobal;  
main () {  
}
```

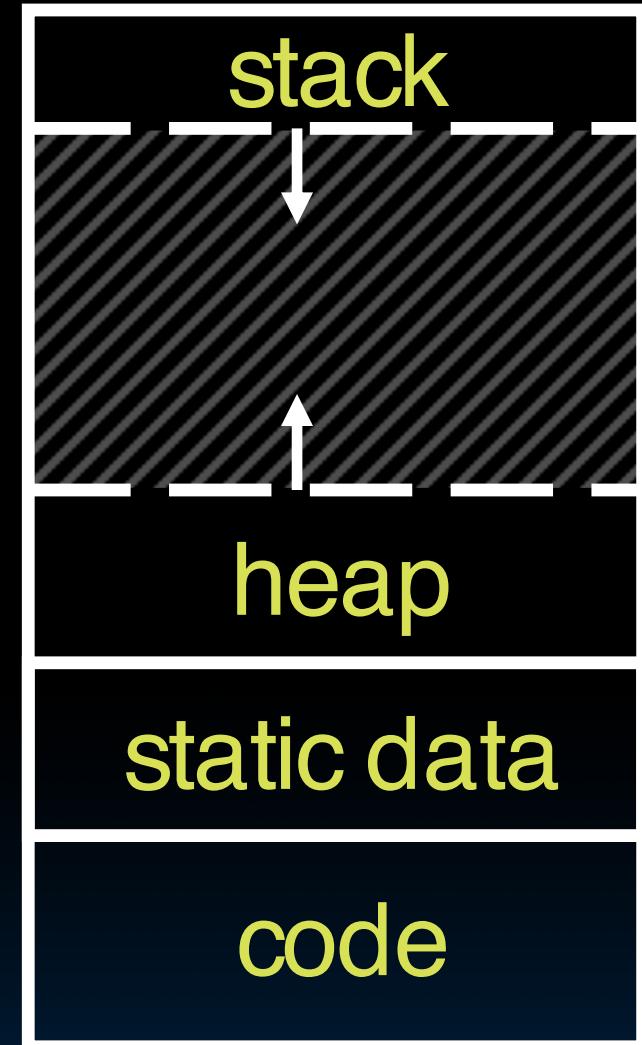
C Memory Management

- **C has 3 pools of memory**
 - **Static storage**: global variable storage, basically permanent, entire program run
 - **The Stack**: local variable storage, parameters, return address (location of “activation records” in Java or “stack frame” in C)
 - **The Heap** (dynamic malloc storage): data lives until deallocated by programmer
- **C requires knowing where objects are in memory, otherwise things don't work as expected**
 - Java hides location of objects

Normal C Memory Management

- A program's **address space** contains 4 regions:
 - **stack**: local variables, grows downward
 - **heap**: space requested for pointers via `malloc()`; resizes dynamically, grows upward
 - **static data**: variables declared outside main, does not grow or shrink
 - **code**: loaded when program starts, does not change

$\sim \text{FFFF}_{\text{hex}}$



$\sim 0_{\text{hex}}$

For now, OS somehow prevents accesses between stack and heap (gray hash lines). Wait for virtual memory

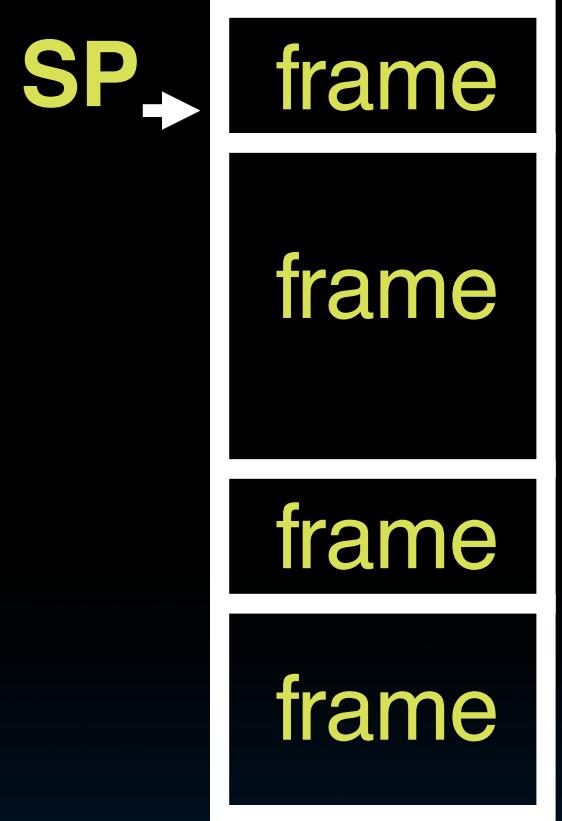
Where are variables allocated?

- If declared **outside** a procedure (global), allocated in “static” storage
- If declared **inside** procedure (local), allocated on the “stack” and **freed when procedure returns.**
 - NB: `main()` is a procedure

```
int myGlobal;  
main () {  
    int myTemp;  
}
```

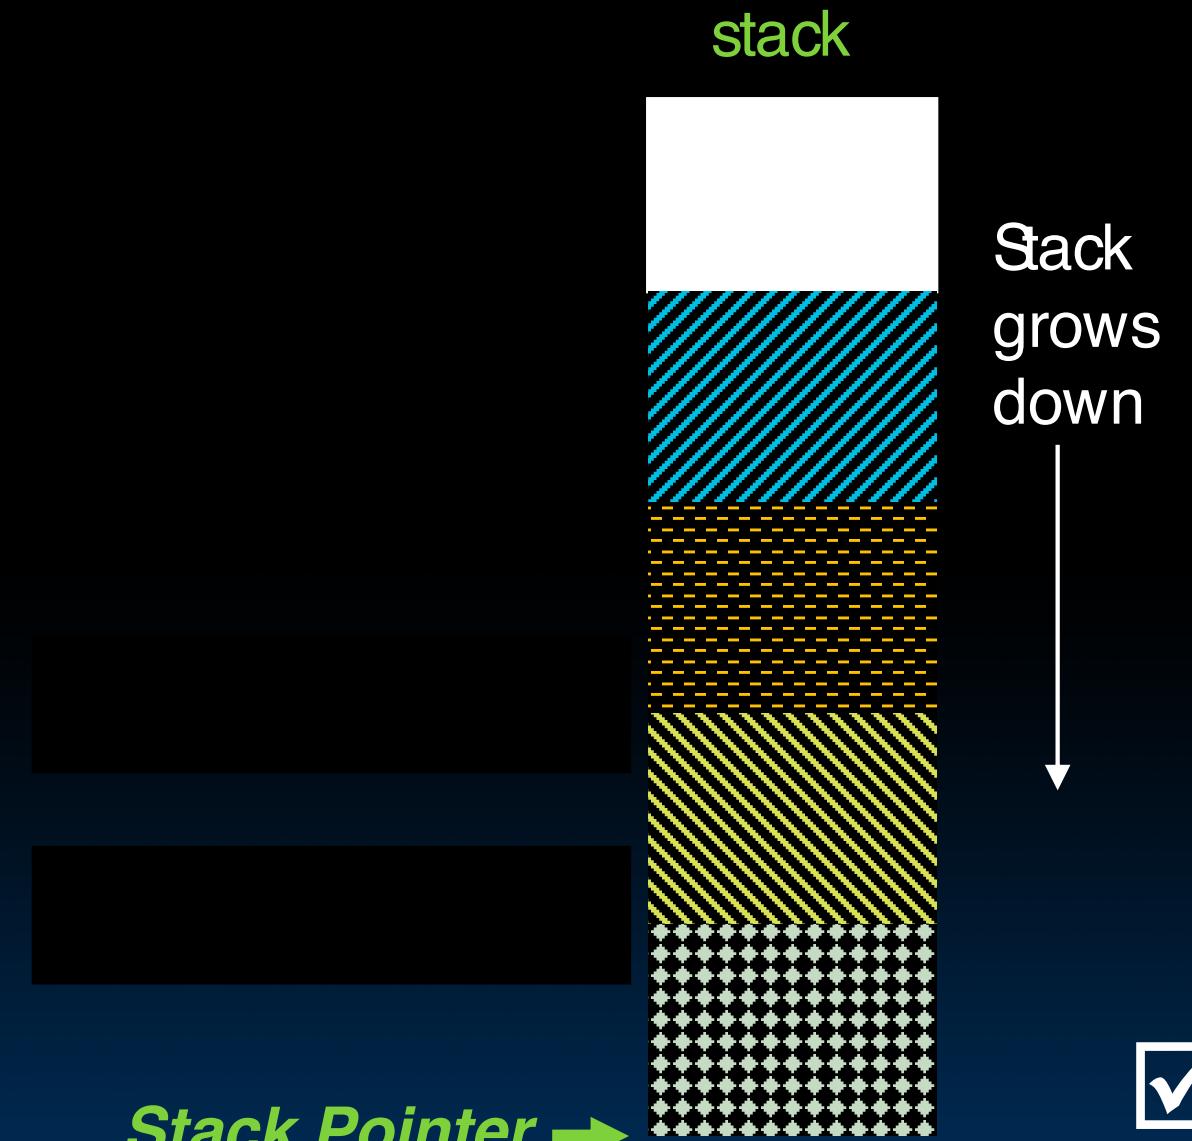
The Stack

- **Stack frame includes:**
 - Return “instruction” address
 - Parameters
 - Space for other local variables
- **Stack frames contiguous blocks of memory; stack pointer tells where top stack frame is**
- **When procedure ends, stack frame is tossed off the stack; frees memory for future stack frames**



- Last In, First Out (LIFO) data structure

```
main ()  
{ a(0);  
}  
void a (int m)  
{ b(1);  
}  
void b (int n)  
{ c(2);  
}  
void c (int o)  
{ d(3);  
}  
void d (int p)  
{  
}
```



Memory Management

The Heap (Dynamic memory)

- Large pool of memory,
not allocated in contiguous order
 - back-to-back requests for heap memory could result blocks very far apart
 - where Java new command allocates memory
- In C, specify number of bytes of memory explicitly to allocate item

```
int *ptr;  
ptr = (int *) malloc(sizeof(int));  
/* malloc returns type (void *),  
so need to cast to right type */
```

- **malloc()**: Allocates raw, uninitialized memory from heap

Memory Management

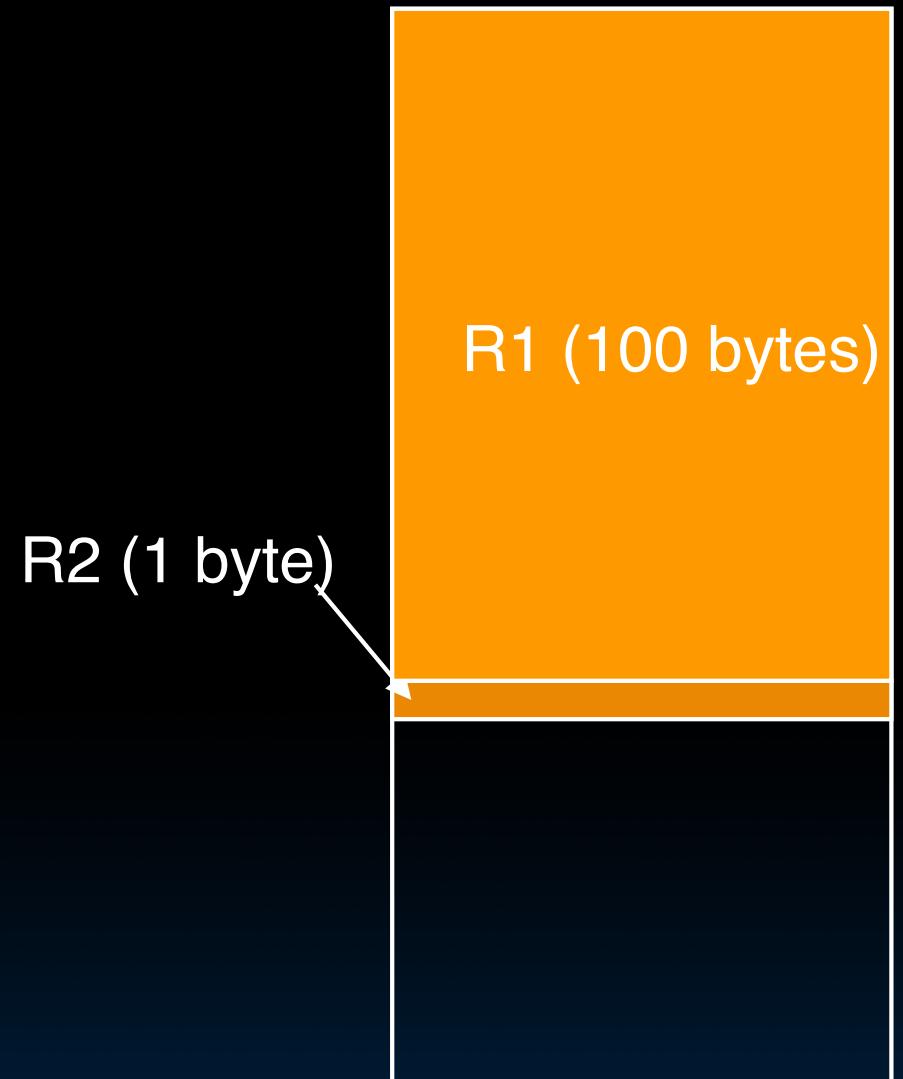
- How do we manage memory?
- Code, Static storage are easy:
 - they never grow or shrink
- Stack space is also easy:
 - stack frames are created and destroyed in last-in, first-out (LIFO) order
- Managing the heap is tricky:
 - memory can be allocated / deallocated at any time

Heap Management Requirements

- Want `malloc()` and `free()` to run quickly
 - Want minimal memory overhead
 - Want to avoid fragmentation* – when most of our free memory is in many small chunks
 - In this case, we might have many free bytes but not be able to satisfy a large request since the free bytes are not contiguous in memory.
- * This is technically called external fragmentation

- An example

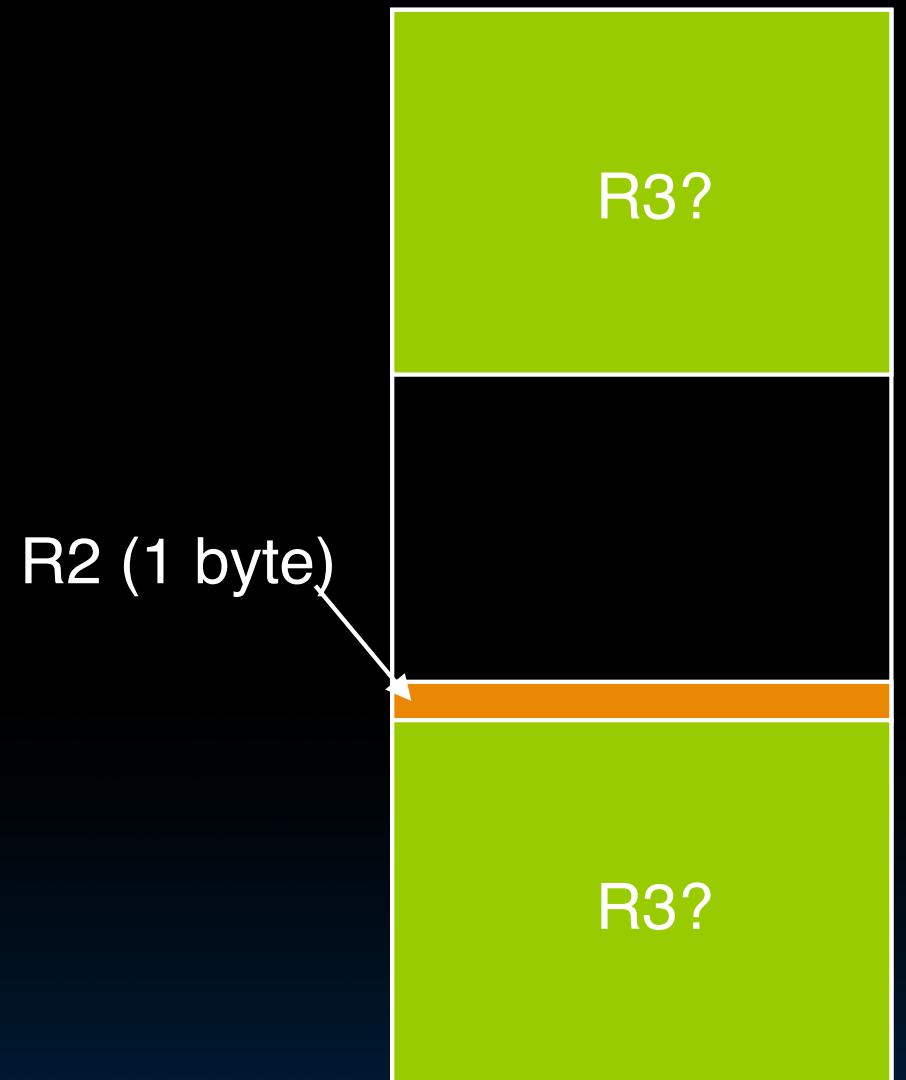
- Request R1 for 100 bytes
- Request R2 for 1 byte
- Memory from R1 is freed
- Request R3 for 50 bytes



Heap Management

- An example

- Request R1 for 100 bytes
- Request R2 for 1 byte
- Memory from R1 is freed
- Request R3 for 50 bytes



K&R Malloc/Free Implementation

- **From Section 8.7 of K&R**
 - Code in the book uses some C language features we haven't discussed and is written in a very terse style, don't worry if you can't decipher the code
- **Each block of memory is preceded by a header that has two fields:**
size of the block and
a pointer to the next block
- **All free blocks are kept in a circular linked list, the pointer field is unused in an allocated block**

K&R Implementation

- **malloc ()** searches the free list for a block that is big enough. If none is found, more memory is requested from the operating system. If what it gets can't satisfy the request, it fails.
- **free ()** checks if the blocks adjacent to the freed block are also free
 - If so, adjacent free blocks are merged (coalesced) into a single, larger free block
 - Otherwise, freed block is just added to the free list

Choosing a block in `malloc()`

- If there are multiple free blocks of memory that are big enough for some request, how do we choose which one to use?
 - best-fit: choose the smallest block that is big enough for the request
 - first-fit: choose the first block we see that is big enough
 - next-fit: like first-fit but remember where we finished searching and resume searching from there

And in conclusion...

- **C has 3 pools of memory**
 - **Static storage**: global variable storage, basically permanent, entire program run
 - **The Stack**: local variable storage, parameters, return address
 - **The Heap (dynamic storage)**: malloc() grabs space from here, free() returns it.
- **malloc () handles free space with freelist**
- **Three ways to find free space when given a request:**
 - **First fit** (find first one that's free)
 - **Next fit** (same as first, but remembers where left off)
 - **Best fit** (finds most "snug" free space)



**When Memory
Goes Bad**

- **Why use pointers?**
 - If we want to pass a huge struct or array, it's easier / faster / etc to pass a pointer than the whole thing.
 - In general, pointers allow cleaner, more compact code.
- **So what are the drawbacks?**
 - Pointers are probably the single largest source of bugs in software, so be careful anytime you deal with them.
 - Dangling reference (use ptr before malloc)
 - Memory leaks (tardy free, lose the ptr)



Writing off the end of arrays...

```
int *foo = (int *) malloc(sizeof(int) * 100);  
int i;  
....  
for(i = 0; i <= 100; ++i) {  
    foo[i] = 0;  
}
```

- **Corrupts other parts of the program...**
 - Including internal C data
- **May cause crashes later**

Returning Pointers into the Stack

- Pointers in C allow access to deallocated memory, leading to hard-to-find bugs !

```
int *ptr () {
    int y;
    y = 3;
    return &y;
}

main () {
    int *stackAddr, content;
    stackAddr = ptr();
    content = *stackAddr;
    printf("%d", content); /* 3 */
    content = *stackAddr;
    printf("%d", content); /*13451514 */
}
```

Use After Free

- When you keep using a pointer..

```
struct foo *f
.....
f = malloc(sizeof(struct foo));
.....
free(f);
.....
bar(f->a);
```

- Reads after the free may be corrupted
 - As something else takes over that memory. Your program will probably get wrong info!
- Writes corrupt other data!
 - Uh oh... Your program crashes later!

Forgetting realloc Can Move Data...

- When you `realloc` it can copy data...
 - ```
struct foo *f = malloc(sizeof(struct foo) * 10);
```

  
...  

```
struct foo *g = f;
```

  
....  

```
f = realloc(sizeof(struct foo) * 20);
```
- Result is `g` *may now point to invalid memory*
  - So reads may be corrupted and writes may corrupt other pieces of memory

# Freeing the Wrong Stuff...

- If you **free()** something never **malloc'ed()**

- Including things like

```
struct foo *f = malloc(sizeof(struct foo) * 10)
```

```
...
```

```
f++;
```

```
...
```

```
free(f)
```

- **malloc** or **free** may get confused..

- Corrupt its internal storage or erase other data...

# Double-Free...

- Eg.,

```
struct foo *f = (struct foo *)
 malloc(sizeof(struct foo) * 10);

...
free(f);

...
free(f);
```

- May cause either a use after free (because something else called malloc() and got that data) or corrupt malloc's data (because you are no longer freeing a pointer called by malloc)

# Losing the initial pointer! (Memory Leak)

- What is wrong with this code?

```
int *plk = NULL;
void genPLK() {
 plk = malloc(2 * sizeof(int));
 ...
 ...
 plk++;
}
```

This MAY be a memory leak  
if we don't keep somewhere else  
a copy of the original malloc'ed  
pointer

# Valgrind to the rescue...

- **Valgrind slows down your program by an order of magnitude, but..**
  - It adds a tons of checks designed to catch most (but not all) memory errors
    - Memory leaks
    - Misuse of free
    - Writing over the end of arrays
- **Tools like Valgrind are absolutely essential for debugging C code**

# And In Conclusion, ...

- **C has three main memory segments in which to allocate data:**
  - Static Data: Variables outside functions
  - Stack: Variables local to function
  - Heap: Objects explicitly malloc-ed/free-d.
- **Heap data is biggest source of bugs in C code**

